



Does prediction improve while the coefficients move?

Two-repetition means. Relative MSE uses the paper's squared-error ratio, without a square root.

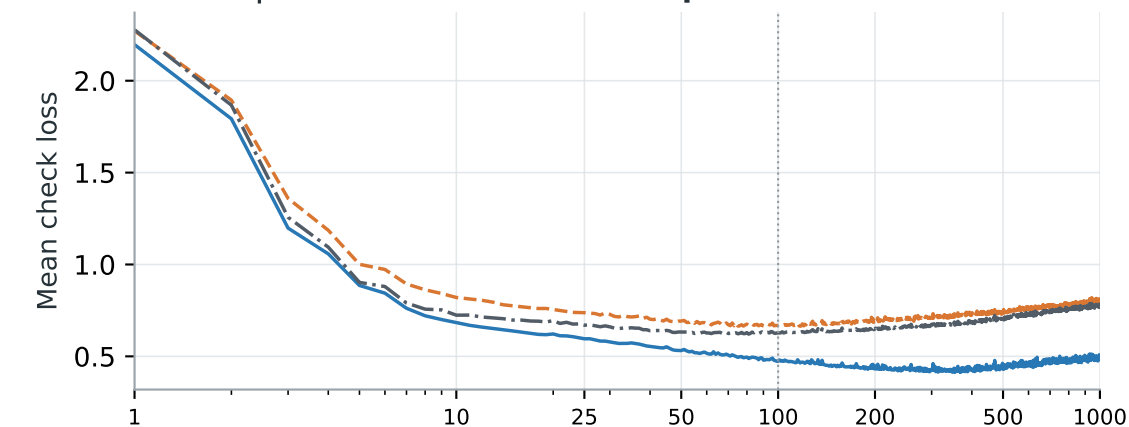
Case 1 | Prediction loss at each epoch



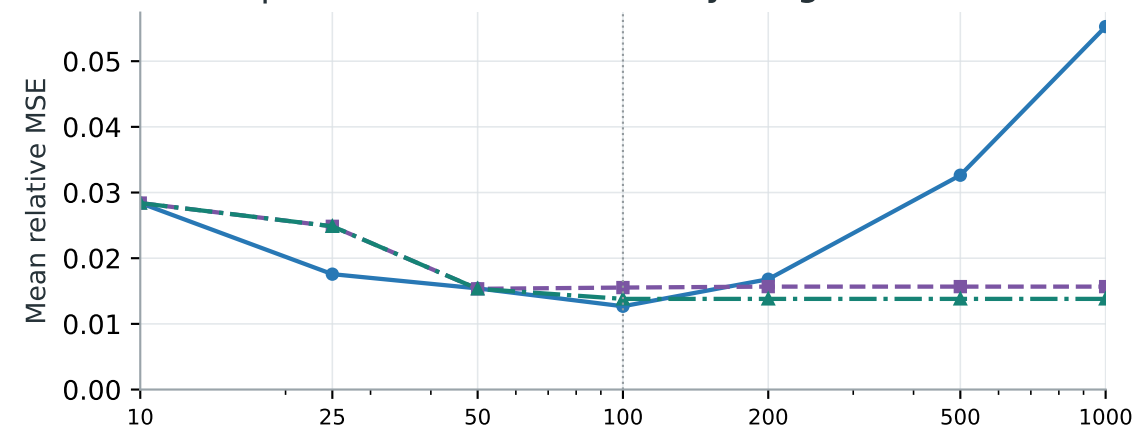
Case 1 | Nuisance-function error by budget



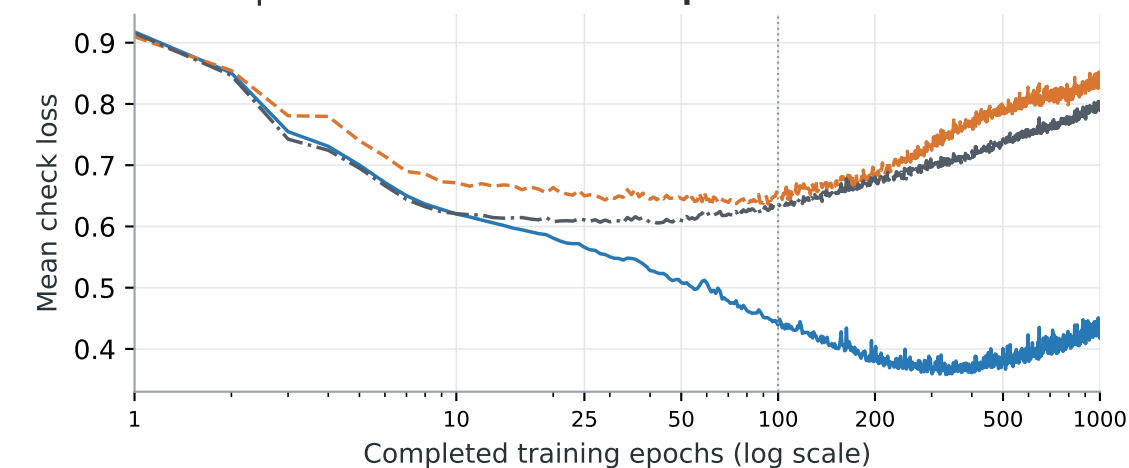
Case 2 | Prediction loss at each epoch



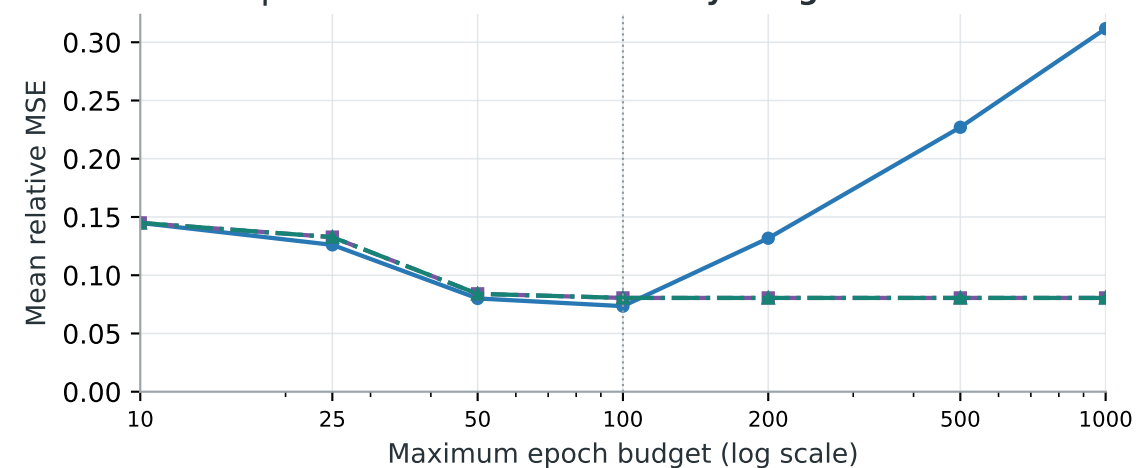
Case 2 | Nuisance-function error by budget



Case 3 | Prediction loss at each epoch



Case 3 | Nuisance-function error by budget



— Training check loss - - Validation check loss - · - Test check loss
—●— At the final epoch - -■- Best validation epoch within budget - -▲- Early stopping: patience 15

Test loss and known truth are diagnostics only; neither is used to select the epoch.