


Do the DPLQR coefficients drift as training continues?

Early stopping is disabled for the paths; the dotted reference retains the first-pass stopping rule.



10 repetitions per case; $n = 1,000$, median quantile. Each path keeps its data and initialization fixed.